

@bind (macro with 1 method)

```

1 begin
2   ### A Pluto.jl notebook ###
3   # v0.19.29
4
5   using Markdown
6   using InteractiveUtils
7
8   # This Pluto notebook uses @bind for interactivity. When running this notebook
9   # outside of Pluto, the following 'mock version' of @bind gives bound variables a
10  # default value (instead of an error).
11  macro bind(def, element)
12    quote
13      local iv = try Base.loaded_modules[Base.PkgId(Base.UUID("6e696c72-6542-
14      2067-7265-42206c756150"), "AbstractPlutoDingetjes").Bonds.initial_value catch; b
15      -> missing; end
16      local el = $(esc(element))
17      global $(esc(def)) = Core.applicable(Base.get, el) ? Base.get(el) :
18      iv(el)
19    end
20  end
21 end

```

Table of Contents

Training a CNN classifier using Flux

Defining the Classifier

Training the network

Testing the network

```

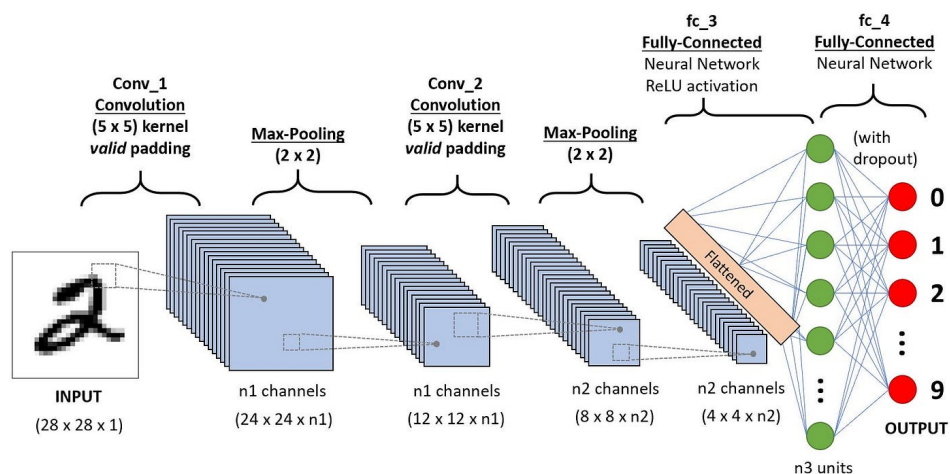
1 begin
2   using PlutoUI
3   using Latexify
4   TableOfContents()
5 end

```

```

1 using Flux: crossentropy, Momentum

```



```

1 md"""
2 ![CNN_1](https://user-images.githubusercontent.com/1801654/224518240-bfe1ea42-3ba4-
3 474c-9e11-e68c1da760be.jpg)
4 """

```

Training a CNN classifier using Flux

MNIST is a dataset of 60k small training images of handwritten digits from 0 to 9.

We will do the following steps in order:

- Load MNIST training and test datasets
- Define a Convolution Neural Network (CNN)
- Define a loss function
- Train the network on the training data
- Test the network on the test data

```
1 md"""
2 ## Training a CNN classifier using `Flux`
3
4 [MNIST](https://www.cs.toronto.edu/~kriz/mnist.html) is a dataset of 60k small
5 training images of handwritten digits from 0 to 9.
6
7 We will do the following steps in order:
8
9 - Load MNIST training and test datasets
10 - Define a Convolution Neural Network (CNN)
11 - Define a loss function
12 - Train the network on the training data
13 - Test the network on the test data
14 """
```

1

```
1 begin
2     using Statistics
3     using CUDA
4     using Flux, Flux.Optimise
5     using MLDatasets: CIFAR10
6     using Images.ImageCore
7     using Flux: onehotbatch, onecold
8     using Base.Iterators: partition
9
10    using Plots
11 end
```

Package cuDNN not found in current path.
 - Run `import Pkg; Pkg.add("cuDNN")` to install the cuDNN package, then restart Julia.
 - If cuDNN is not installed, some Flux functionalities will not be available when running on the GPU.

Tip

The preceding cell needs to be modified to read in the CIFAR10 database rather than MNIST.

```
1 md"""
2 !!! tip
3     The preceding cell needs to be modified to read in the CIFAR10 database rather
4     than MNIST.
5 """
```

Set an environment variable to stop the system asking for approval to download data.

```
1 md"""
2 Set an environment variable to stop the system asking for approval to download data.
3 """
```

```
1 ENV["DATADEPS_ALWAYS_ACCEPT"] = true;
```

```
1 begin
2     train_x, train_y = CIFAR10(split=:train)[: ]
3     train_x = reshape(train_x, 32, 32, 3, :)
4     labels2 = onehotbatch(train_y, 0:9)
5 end
```

```
Extracting archive:
--
Path =
Type = tar
Code Page = UTF-8
Characteristics = ASCII
Everything is Ok

Folders: 1
Files: 8
Size: 184380149
Compressed: 9216
```

```
1 md"""
2 The train\_x array contains 60,000 28x28 pixel images converted to 28 x 28 x 1 arrays
  with the third dimension acting as a grey-scale channel. Let's take a look at a
  random image from train\_x. For this, we need to convert the floats in the 28x28
  matrix to the grey colour value:
3 """
```

The size and dimensionality of the CIFAR10 data base are different: there are 50,000 images of 32 x 32 pixels involving 3 colour channels

```
1 md"""
2 !!! tip
3     The size and dimensionality of the CIFAR10 data base are different: there are
4     50,000 images of 32 x 32 pixels involving 3 colour channels
5 """
```

```

3×32×32 Array{Gray{Float32},3} with eltype Gray{Float32}:
[:, :, 1] =
Gray{Float32}(0.917647) Gray{Float32}(0.556863) ... Gray{Float32}(0.619608)
Gray{Float32}(0.92549) Gray{Float32}(0.552941) Gray{Float32}(0.513726)
Gray{Float32}(0.901961) Gray{Float32}(0.521569) Gray{Float32}(0.411765)

[:, :, 2] =
Gray{Float32}(0.839216) Gray{Float32}(0.435294) ... Gray{Float32}(0.607843)
Gray{Float32}(0.843137) Gray{Float32}(0.435294) Gray{Float32}(0.498039)
Gray{Float32}(0.831373) Gray{Float32}(0.415686) Gray{Float32}(0.439216)

[:, :, 3] =
Gray{Float32}(0.717647) Gray{Float32}(0.45098) ... Gray{Float32}(0.533333)
Gray{Float32}(0.713726) Gray{Float32}(0.45098) Gray{Float32}(0.411765)
Gray{Float32}(0.717647) Gray{Float32}(0.443137) Gray{Float32}(0.345098)

...

[:, :, 30] =
Gray{Float32}(0.372549) Gray{Float32}(0.368627) ... Gray{Float32}(0.772549)
Gray{Float32}(0.309804) Gray{Float32}(0.294118) Gray{Float32}(0.788235)
Gray{Float32}(0.160784) Gray{Float32}(0.152941) Gray{Float32}(0.690196)

[:, :, 31] =
Gray{Float32}(0.490196) Gray{Float32}(0.501961) ... Gray{Float32}(0.701961)
Gray{Float32}(0.407843) Gray{Float32}(0.396078) Gray{Float32}(0.701961)
Gray{Float32}(0.294118) Gray{Float32}(0.286275) Gray{Float32}(0.611765)

[:, :, 32] =
Gray{Float32}(0.780392) Gray{Float32}(0.627451) ... Gray{Float32}(0.545098)
Gray{Float32}(0.709804) Gray{Float32}(0.52549) Gray{Float32}(0.556863)
Gray{Float32}(0.635294) Gray{Float32}(0.45098) Gray{Float32}(0.462745)

1 Gray.(permutedims(reshape(train_x[:, :, :, rand(1:end)], 32, 32, 3), (3, 2, 1)))
2

```

Tip

Use this function: `image(x) = colorview(Gray, permutedims(x, (3, 2, 1)))` to display the colour images instead of Gray.

```

1 md"""
2 !!! tip
3     Use this function: `image(x) = colorview(Gray, permutedims(x, (3, 2, 1)))` to
4     display the colour images instead of `Gray`.
5 """

```

The images are simply 28 x 28 matrices of numbers plus one greyscale channel. We can now arrange them in batches of say, 1,000 and keep a validation set to track our progress. This process is called minibatch learning, which is a popular method of training large neural networks. Rather than sending the entire dataset at once, we break it down into smaller chunks (called minibatches) typically chosen at random.

The first 59k images (in batches of 1,000) will be our training set, and the rest are for validation used to track training progress. `partition` handily breaks down the set we give it in consecutive parts (1,000 in this case).

```

1 md"""
2 The images are simply 28 x 28 matrices of numbers plus one greyscale channel. We can
3 now arrange them in batches of say, 1,000 and keep a validation set to track our
4 progress. This process is called minibatch learning, which is a popular method of
5 training large neural networks. Rather than sending the entire dataset at once, we
6 break it down into smaller chunks (called minibatches) typically chosen at random.
7
8 The first 59k images (in batches of 1,000) will be our training set, and the rest are
9 for validation used to track training progress. `partition` handily breaks down the
10 set we give it in consecutive parts (1,000 in this case).
11 """

```



```

1 begin
2
3     # Function to convert images for display
4     image(x) = colorview(RGB, permutedims(x, (3, 2, 1)))
5
6     # Organize into minibatches (each containing 1000 images)
7     batch_size = 1000
8     num_batches = div(size(train_x, 4), batch_size)
9
10    # Split into training and validation sets
11    train_set = [train_x[:, :, :, (i-1)*batch_size+1:i*batch_size] for i in
12    1:num_batches-1]
13    validation_set = train_x[:, :, :, end-(batch_size-1):end]
14
15    # Display the first image in the validation set as an example
16    image(validation_set[:, :, :, 20])
17 end

```

Tip

Because CIFAR10 comprises 50,000 images rather than 60,000, the split between training and validation needs to reflect this.

```

1 md"""
2 !!! tip
3     Because CIFAR10 comprises 50,000 images rather than 60,000, the split between
4     training and validation needs to reflect this.
5 """

```

10×50000 OneHotMatrix{::Vector{UInt32}} with eltype Bool:

```

. . . . . 1 . . . . . 1 . . . . . ... . . . . . 1 . . . . . 1 . . . . .
. 1 . . 1 . . . . . . . . . . 1 . 1 . . . . 1 . . . . 1 . . . . 1 1
. . . . . 1 1 . . . . . . . . . . 1 . . . . 1 . . . . 1 . . . . .
. . . . . . . . . . . . . . . 1 . . . . . . . . . . . . . . . .
. . 1 . . . . . . . . . . . . . . . . . . . 1 . . . . . . . . . .
. . . . . . . . . . 1 . . . . . . . . . . ... . 1 . . . . . . . . . .
. . . . . . . . . . . 1 . . . . . . . . . . . . . . . . . . . 1 . . . .
1 . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . .
. . . . . . . . . . . . . . . . . . . . . . . . . . . . . . . .
. . . . . . . . . . 1 . . . . . . . . . . . . . . . . . . . 1 . . .

```

```

1 begin
2     train = [(train_x[:, :, :, i], labels2[:, i]) for i in partition(1:50000, 1000)]
3     |> gpu
4     valset = 45001:50000
5     valX = train_x[:, :, :, valset] |> gpu
6     valY = labels2[:, valset] |> gpu
7 end

```

The CUDA function is being called but CUDA.jl is not functional.
Defaulting back to the CPU. (No action is required if you want to run on the CPU).

Defining the Classifier

Now we can define our Convolutional Neural Network (CNN).

A convolutional neural network is one which defines a kernel and slides it across a matrix to create an intermediate representation to extract features from. It creates higher order features as it goes into deeper layers, making it suitable for images, where the structure of the subject is what will help us determine which class it belongs to.

```
1 md"""
2 ### Defining the Classifier
3
4 Now we can define our Convolutional Neural Network (CNN).
5
6 A convolutional neural network is one which defines a kernel and slides it across a
  matrix to create an intermediate representation to extract features from. It creates
  higher order features as it goes into deeper layers, making it suitable for images,
  where the structure of the subject is what will help us determine which class it
  belongs to.
7 """
```

```
m3 = Chain(
  Conv((5, 5), 3 => 32, relu, pad=2), # 2_432 parameters
  MaxPool((2, 2)),
  Conv((5, 5), 32 => 16, relu, pad=2), # 12_816 parameters
  MaxPool((2, 2)),
  Main.var"#11#12"{typeof(reshape), typeof(size), Colon}(reshape, size, Colon()),
  Dense(1024 => 120), # 123_000 parameters
  Dense(120 => 84), # 10_164 parameters
  Dense(84 => 10), # 850 parameters
  NNlib.softmax,
) # Total: 10 arrays, 149_262 parameters, 584.289 KiB.
```

```
1 # Define the neural network architecture
2 m3 = Chain(
3   Conv((5, 5), 3 => 32, pad=(2, 2), relu), # Modify input filter number from 1 to
  3 for 3 input channels
4   MaxPool((2, 2)),
5   Conv((5, 5), 32 => 16, pad=(2, 2), relu), # Modify output channels from 16 to 32
6   MaxPool((2, 2)),
7   x -> reshape(x, :, size(x, 4)),
8   Dense(1024, 120), # Modify input dimension from 392 to 1024 for the first Dense
  layer
9   Dense(120, 84),
10  Dense(84, 10),
11  softmax
12 ) |> gpu
```

Tip

You need to modify the input filter number of the first Conv layer and the input dimension of the first Dense layer.

```
1 md"""
2 !!! tip
3   You need to modify the input filter number of the first 'Conv' layer and the
  input dimension of the first 'Dense' layer.
4 """
```

We will use a crossentropy loss and the Momentum optimiser here. Crossentropy will be a good option when it comes to working with multiple independent classes. Momentum smooths out the noisy gradients and helps towards a smooth convergence. Gradually lowering the learning rate along with momentum helps to maintain a bit of adaptivity in our optimisation, preventing us from overshooting our desired destination.

```
1 md"""
2 We will use a crossentropy loss and the Momentum optimiser here. Crossentropy will be
  a good option when it comes to working with multiple independent classes. Momentum
  smooths out the noisy gradients and helps towards a smooth convergence. Gradually
  lowering the learning rate along with momentum helps to maintain a bit of adaptivity
  in our optimisation, preventing us from overshooting our desired destination.
3 """
```

```
Momentum(0.01, 0.9, IdDict())
```

```
1 begin
2   loss2(x, y) = sum(crossentropy(m3(x), y))
3   opt2 = Momentum(0.01)
4 end
```

We can start writing our train loop where we will keep track of some basic accuracy numbers about our model. We can define an accuracy function for it like so:

```
1 md"""
2 We can start writing our train loop where we will keep track of some basic accuracy
3 numbers about our model. We can define an 'accuracy' function for it like so:
4 """
```

accuracy (generic function with 1 method)


```
1 accuracy(x, y) = mean(onecold(m3(x), 0:9) .== onecold(y, 0:9))
```

Training the network

Training is where we do a bunch of the interesting operations we defined earlier, and see what our net is capable of. We will loop over the dataset 10 times and feed the inputs to the neural network and optimise.

```
1 md"""
2 ### Training the network
3
4 Training is where we do a bunch of the interesting operations we defined earlier, and
5 see what our net is capable of. We will loop over the dataset 10 times and feed the
6 inputs to the neural network and optimise.
7 """
```

```
1 begin
2   train_acc = Float32[]
3   train_epochs = Int32[]
4   epochs = 10
5
6   for epoch = 1:epochs
7     for d in train
8       gs = gradient(Flux.params(m3)) do
9         l = loss2(d...)
10      end
11      update!(opt2, Flux.params(m3), gs)
12    end
13    push!(train_acc, accuracy(valX, valY))
14    push!(train_epochs, epoch)
15  end
16
17 end
```



```

1 begin
2     plot(train_epochs, train_acc, lab="")
3     yaxis!("Accuracy")
4     xaxis!("Training epoch")
5 end

```

Tip

The training regime doesn't need modification.

```

1 md"""
2 !!! tip
3     The training regime doesn't need modification.
4 """

```

Testing the network

We have trained the network for 10 passes over the training dataset. But we need to check if the network has learnt anything at all.

We will check this by predicting the class label that the neural network outputs, and checking it against the ground-truth. If the prediction is correct, we add the sample to the list of correct predictions. This will be done on a yet unseen section of data.

First step. Let us perform the exact same preprocessing on this set, as we did on our training set.

```

1 md"""
2 ### Testing the network
3
4 We have trained the network for 10 passes over the training dataset. But we need to
  check if the network has learnt anything at all.
5
6 We will check this by predicting the class label that the neural network outputs, and
  checking it against the ground-truth. If the prediction is correct, we add the sample
  to the list of correct predictions. This will be done on a yet unseen section of data.
7
8 First step. Let us perform the exact same preprocessing on this set, as we did on our
  training set.
9
10 """

```

Next, display images from the test set.

```

1 md"""
2 Next, display images from the test set.
3 """

```



```
32×32×3×10000 Array{Float32, 4}:
[:, :, 1, 1] =
0.619608 0.596078 0.592157 0.607843 0.607843 ... 0.266667 0.239216 0.211765
0.623529 0.592157 0.592157 0.607843 0.611765 0.164706 0.192157 0.219608
0.647059 0.623529 0.619608 0.627451 0.631373 0.121569 0.137255 0.176471
0.65098 0.65098 0.654902 0.682353 0.666667 0.14902 0.168627 0.168627
0.627451 0.635294 0.627451 0.654902 0.662745 0.145098 0.152941 0.156863
0.611765 0.627451 0.639216 0.654902 0.639216 ... 0.168627 0.164706 0.156863
0.635294 0.643137 0.647059 0.662745 0.662745 0.164706 0.172549 0.156863
⋮
0.54902 0.568627 0.560784 0.443137 0.333333 0.309804 0.427451 0.411765
0.552941 0.556863 0.54902 0.517647 0.411765 0.286275 0.364706 0.34902
0.560784 0.560784 0.556863 0.54902 0.501961 0.219608 0.235294 0.188235
0.537255 0.533333 0.545098 0.54902 0.541176 0.14902 0.101961 0.0941176
0.494118 0.490196 0.509804 0.533333 0.521569 ... 0.0509804 0.113725 0.133333
0.454902 0.466667 0.470588 0.498039 0.505882 0.156863 0.0784314 0.0823529
```

```
[:, :, 2, 1] =
0.439216 0.439216 0.431373 0.419608 ... 0.423529 0.486275 0.454902 0.419608
0.435294 0.431373 0.427451 0.431373 0.380392 0.392157 0.4 0.411765
0.454902 0.447059 0.435294 0.427451 0.360784 0.345098 0.333333 0.34902
0.462745 0.454902 0.435294 0.439216 0.345098 0.356863 0.356863 0.337255
0.439216 0.439216 0.415686 0.431373 0.345098 0.341176 0.352941 0.34902
0.427451 0.443137 0.45098 0.458824 ... 0.329412 0.34902 0.360784 0.360784
0.45098 0.458824 0.458824 0.470588 0.521569 0.309804 0.345098 0.341176
⋮
0.384314 0.4 0.403922 0.333333 0.4 0.517647 0.611765 0.572549
0.380392 0.380392 0.388235 0.384314 0.490196 0.513726 0.568627 0.529412
0.380392 0.384314 0.388235 0.4 0.431373 0.454902 0.45098 0.388235
0.372549 0.372549 0.384314 0.396078 0.352941 0.380392 0.321569 0.301961
0.356863 0.356863 0.372549 0.388235 ... 0.235294 0.25098 0.321569 0.329412
0.333333 0.345098 0.34902 0.368627 0.364706 0.333333 0.25098 0.262745
```

```
1 begin
2     test_x, test_y = CIFAR10(split=:test)[: ]
3     test_x = reshape(test_x, 32, 32, 3, :)
4 end
```

```
1 test_labels = onehotbatch(test_y, 0:9);
```

```
test =
[(32×32×3×1000 Array{Float32, 4}:
[:, :, 1, 1] =
```

```
1 test = ([test_x[:, :, i], test_labels[:, i]] for i in partition(1:10000, 1000)) |>
gpu
```

```
1 @bind image_index Slider(1:100:10000, default=1)
```

```
3×32×32 Array{Gray{Float32}, 3} with eltype Gray{Float32}:
[:, :, 1] =
Gray{Float32}(0.619608) Gray{Float32}(0.596078) ... Gray{Float32}(0.211765)
Gray{Float32}(0.439216) Gray{Float32}(0.439216) Gray{Float32}(0.419608)
Gray{Float32}(0.192157) Gray{Float32}(0.2) Gray{Float32}(0.627451)

[:, :, 2] =
Gray{Float32}(0.623529) Gray{Float32}(0.592157) ... Gray{Float32}(0.219608)
Gray{Float32}(0.435294) Gray{Float32}(0.431373) Gray{Float32}(0.411765)
Gray{Float32}(0.184314) Gray{Float32}(0.156863) Gray{Float32}(0.584314)

[:, :, 3] =
Gray{Float32}(0.647059) Gray{Float32}(0.623529) ... Gray{Float32}(0.176471)
Gray{Float32}(0.454902) Gray{Float32}(0.447059) Gray{Float32}(0.34902)
Gray{Float32}(0.2) Gray{Float32}(0.176471) Gray{Float32}(0.517647)

⋮

[:, :, 30] =
Gray{Float32}(0.537255) Gray{Float32}(0.533333) ... Gray{Float32}(0.0941176)
Gray{Float32}(0.372549) Gray{Float32}(0.372549) Gray{Float32}(0.301961)
Gray{Float32}(0.141176) Gray{Float32}(0.121569) Gray{Float32}(0.486275)

[:, :, 31] =
Gray{Float32}(0.494118) Gray{Float32}(0.490196) ... Gray{Float32}(0.133333)
Gray{Float32}(0.356863) Gray{Float32}(0.356863) Gray{Float32}(0.329412)
Gray{Float32}(0.141176) Gray{Float32}(0.12549) Gray{Float32}(0.505882)

[:, :, 32] =
Gray{Float32}(0.454902) Gray{Float32}(0.466667) ... Gray{Float32}(0.0823529)
Gray{Float32}(0.333333) Gray{Float32}(0.345098) Gray{Float32}(0.262745)
Gray{Float32}(0.129412) Gray{Float32}(0.133333) Gray{Float32}(0.431373)
```

```
1 Gray.(permutedims(reshape(test_x[:, :, :, image_index], 32, 32, 3), (3, 2, 1)))
```

The outputs are energies for the 10 classes. Higher the energy for a class, the more the network thinks that the image is of the particular class. Every column corresponds to the output of one image, with the 10 floats in the column being the energies.

Let's see how the model fared:

```
1 md"""
2 The outputs are energies for the 10 classes. Higher the energy for a class, the more
  the network thinks that the image is of the particular class. Every column
  corresponds to the output of one image, with the 10 floats in the column being the
  energies.
3
4 Let's see how the model fared:
5 """
```

```
10x5 Matrix{Float32}:
0.12386      0.160771      0.00802864      0.0013145      1.63209f-5
6.27258f-5    0.0413777      0.00262401      0.000331806      3.84491f-6
0.693682      0.00954001      0.0455476      0.0269539      0.0096018
0.021898      0.0096466      0.167972      0.233842      0.144129
0.0388883     0.00762291      0.119493      0.0296499      0.00107298
0.0200257     0.00188005      0.127114      0.565769      0.834886
0.000224681   0.00941081      0.428894      0.0660704      0.00732907
0.083864      0.00113594      0.0326832      0.0698819      0.00280342
0.0085742     0.66651         0.0641139      0.00232622      1.54209f-5
0.00891988    0.0921051      0.00352822      0.00386034      0.00014131
```

```
1 begin
2     ids = rand(1:10000, 5)
3     rand_test = test_x[:, :, ids] |> gpu
4     rand_truth = test_y[ids]
5     m3(rand_test)
6 end
```

Tip

For visualisation you should display performance on a random sample of test images, say 20, and set up a slider to navigate them while displaying the image, the ground truth label, and the label predicted by the network.

```
1 md"""
2 !!! tip
3 For visualisation you should display performance on a random sample of test
  images, say 20, and set up a slider to navigate them while displaying the image,
  the ground truth label, and the label predicted by the network.
4 """
```

This looks similar to how we would expect the results to be. At this point, it's a good idea to see how our net actually performs on new data, that we have prepared.

```
1 md"""
2 This looks similar to how we would expect the results to be. At this point, it's a
  good idea to see how our net actually performs on new data, that we have prepared.
3 """
```

0.547

```
1 accuracy(test[1]...)
```

This is much better than random chance set at 10% (since we only have 10 classes), and not bad at all for a relatively small hand-coded network like this one.

```
1 md"""
2 This is much better than random chance set at 10% (since we only have 10 classes),
  and not bad at all for a relatively small hand-coded network like this one.
3 """
```

Let's take a look at how the net performed on all the classes performed individually.

```
1 md"""
2 Let's take a look at how the net performed on all the classes performed individually.
3 """
```

```

1 begin
2   class_correct = zeros(10)
3   class_total = zeros(10)
4   for i in 1:10
5     preds = m3(test[i][1])
6     lab = test[i][2]
7     for j = 1:1000
8       pred_class = findmax(preds[:, j])[2]
9       actual_class = findmax(lab[:, j])[2]
10      if pred_class == actual_class
11        class_correct[pred_class] += 1
12      end
13      class_total[actual_class] += 1
14    end
15  end
16 end

```

	accuracy	label
1	0.532	0
2	0.648	1
3	0.37	2
4	0.39	3
5	0.373	4
6	0.52	5
7	0.752	6
8	0.656	7
9	0.728	8
10	0.653	9

```

1 begin
2   using DataFrames
3   DataFrame(accuracy=(class_correct ./ class_total), label=0:9)
4 end

```

Tip

For legibility you should assign labels to the image categories.

```

1 md"""
2 !!! tip
3   For legibility you should assign labels to the image categories.
4 """

```

The spread seems pretty good, with certain classes performing significantly better than the others.

```

1 md"""
2 The spread seems pretty good, with certain classes performing significantly better
3   than the others.
4 """

```

ENTS =
8ef-5323-5732-b1bb-66c8b64840ba\"\\nDataFrames = \"a93c6f00-e57d-5684-b7b6-d8193f3e46c0\"\\nFl

PLUTO_MANIFEST_TOML_CONTENTS =
"# This file is machine-generated - editing it directly is not advised\\n\\njlulia_version = '

